tickerplant.ai

Single Console Tickerplant : q/Kdb+/pyQ/python Stack

Rohan Shiloh Shah
rohanshilohshah@outlook.com

June 2, 2018

www.cs.McGill.ca/~rshah3

Table of contents

Outline of Presentation

'Dispel any notion of tick being a black box product which cannot be modified according(ly)' -N.Perrem

Sections

- What is tick data?

Outline of Presentation

'Dispel any notion of tick being a black box product which cannot be modified according(ly)' -N.Perrem

Sections

- What is tick data?
- Why do traditional RDBMS fail spectacularly when trying to capture this data in real-time for subsequent analysis?

Outline of Presentation

'Dispel any notion of tick being a black box product which cannot be modified according(ly)' -N.Perrem

Sections

- What is tick data?
- Why do traditional RDBMS fail spectacularly when trying to capture this data in real-time for subsequent analysis?
- Overview of q/Kdb+ tickerplant architecture

Outline of Presentation

'Dispel any notion of tick being a black box product which cannot be modified according(ly)' -N.Perrem

Sections

- What is tick data?
- Why do traditional RDBMS fail spectacularly when trying to capture this data in real-time for subsequent analysis?
- Overview of q/Kdb+ tickerplant architecture
- Machine Learning and Order Book Analysis

Temporal Big Data

Order Book

The real-time fx/equity/derivative order book is generally available at one of 4 levels.

- Level 1 : best bid/ask prices with size and traded volume
- Level 2 : market depth 5
- Level 3 : market depth 20
- Tick by Tick : yeverything

Bombay & National Stock Exchanges

The NSE provides feeds for all 4 levels whereas the BSE only provides 1 minute snapshots of the first 2 levels? Derivatives data from the BSE is currently free.

Pricing

`www.nseindia.com/supra_global/content/dotex/data_products.htm`

`www.bseindia.com/products_and_services/productinformation.aspx`

Order Book Size on Disk

Files : Date Partitioned Column Database & q Code

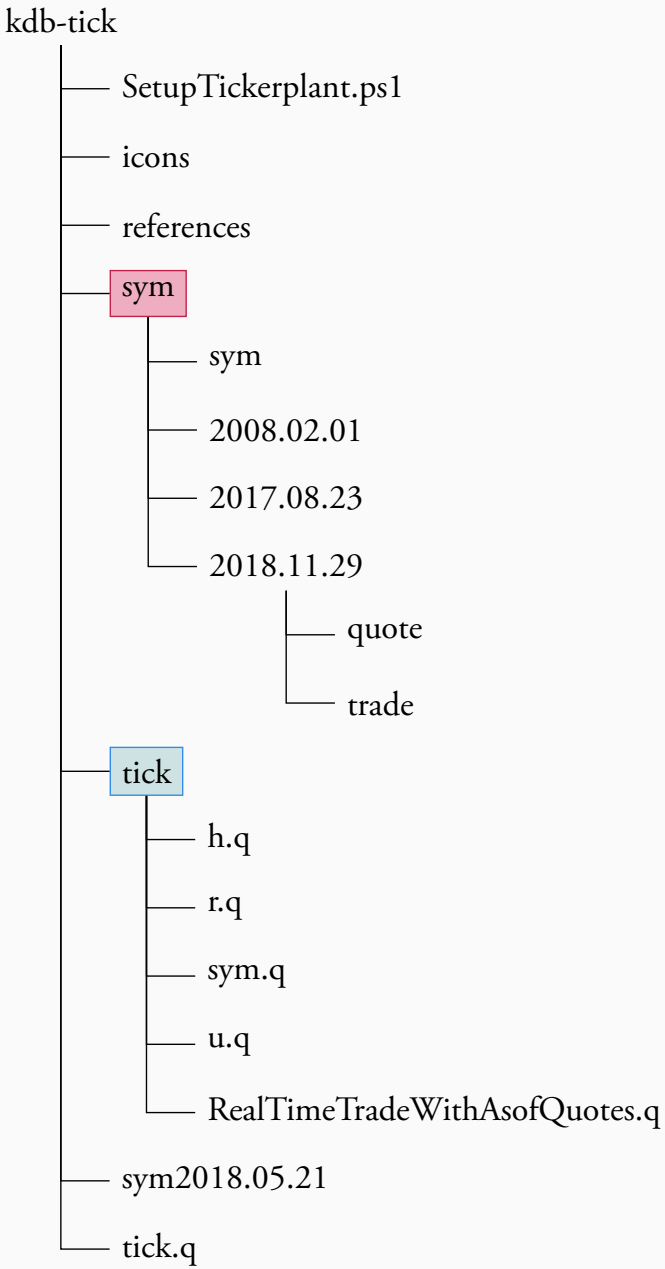

Tables : Quote & Trade

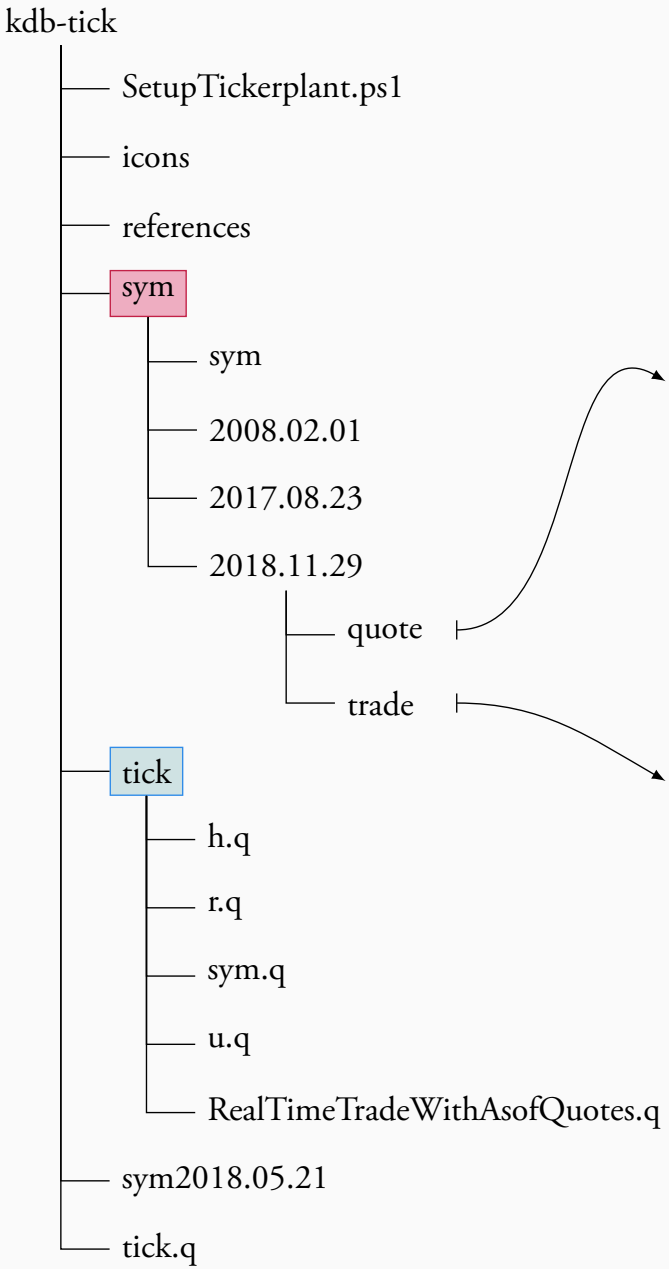

c	t	f	a
date	d		
time	n		(s)
sym	s		g
bid	f		
ask	f		
bsize	j		
asize	j		

c	t	f	a
date	d		
time	n		(s)
sym	s		g
price	f		
size	j		

Joins: Latest Trade

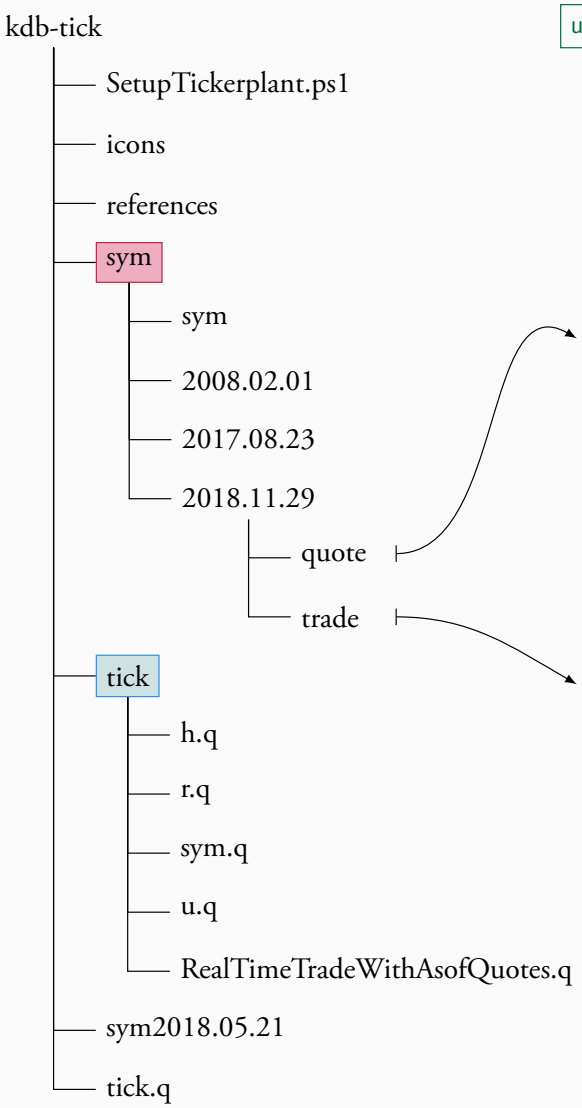


```
upsert select (last) by sym from quote
```

c	t	f	a
date	d		
time	n		(s)
sym	s		g
bid	f		
ask	f		
bsize	j		
asize	j		

c	t	f	a
date	d		
time	n		(s)
sym	s		g
price	f		
size	j		

Sym	Time	Bid	Ask	Bsize	Asize
IBM.N	0D08:04:45.262855000	191.17	191.19	422	572
MSFT.O	0D08:04:45.162728000	45.15	45.16	609	884

Temporal Joins : Latest Bid/Ask Asof Trade : Method 1

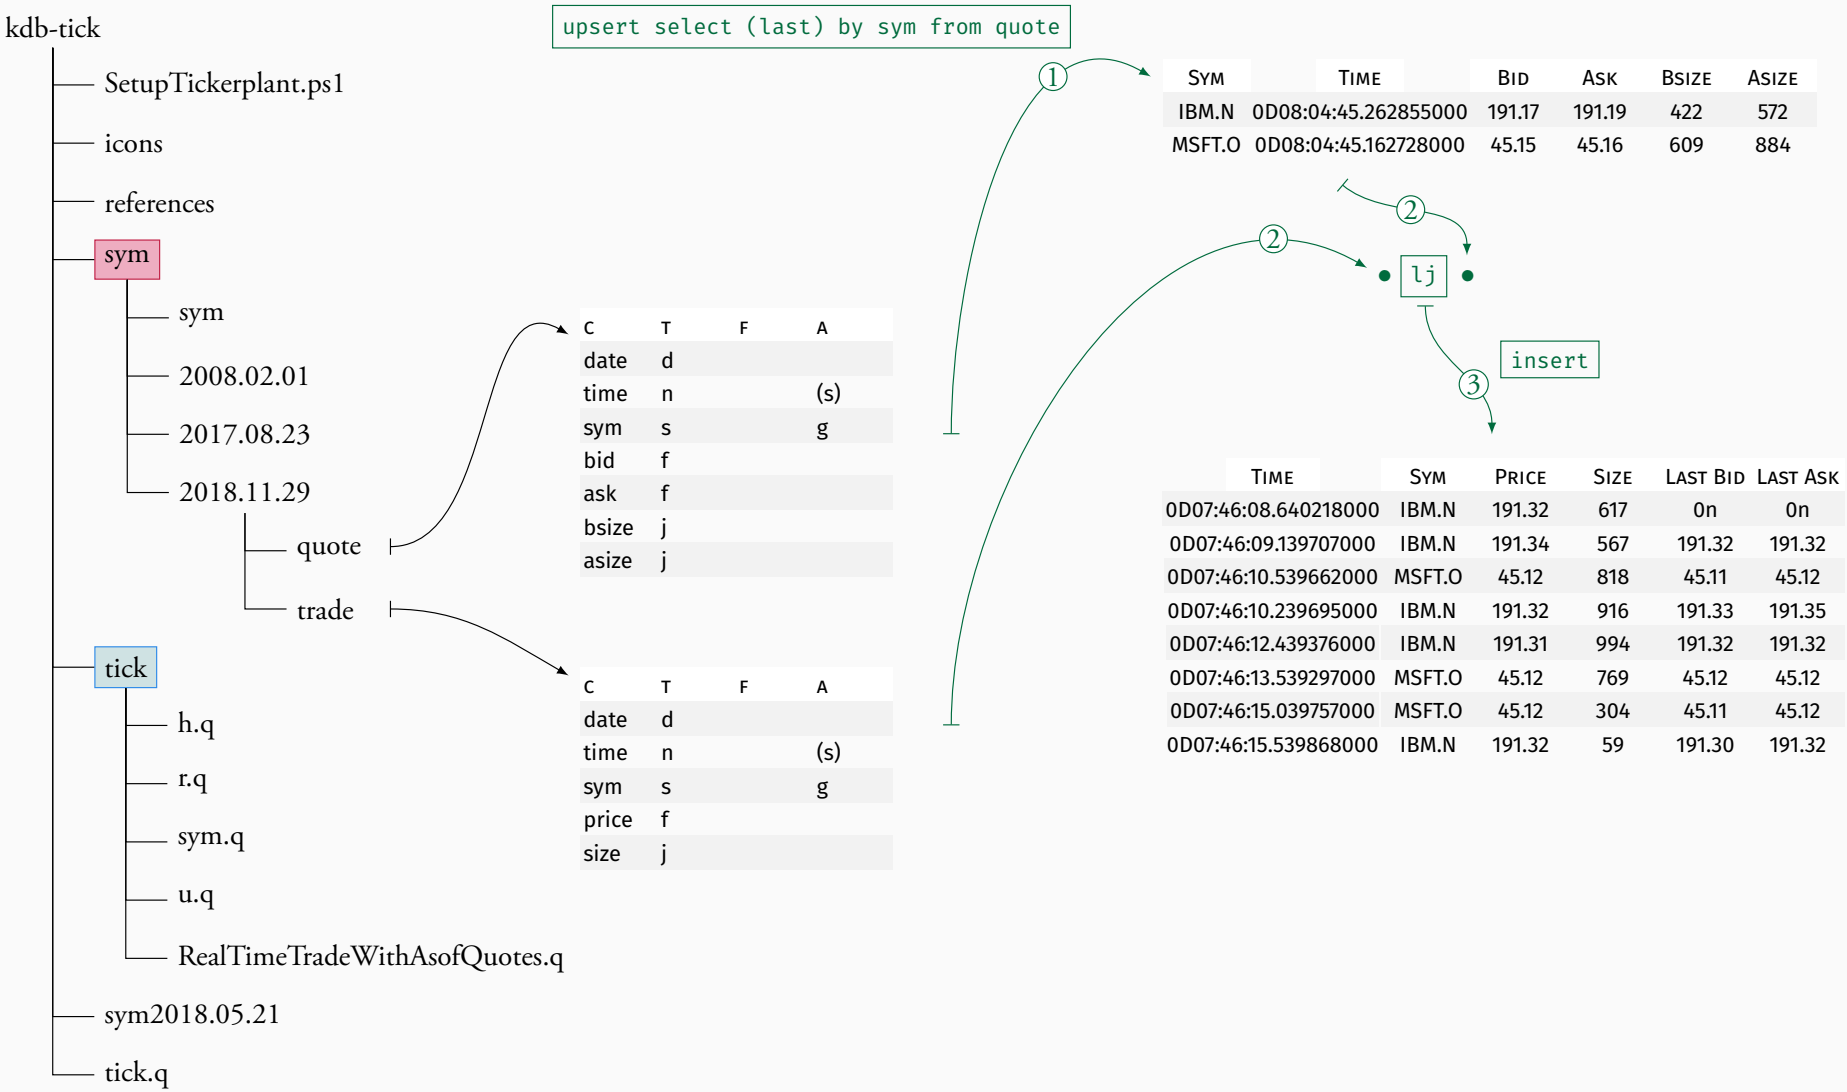

Temporal Joins : Latest Bid/Ask Asof Trade : Method 2

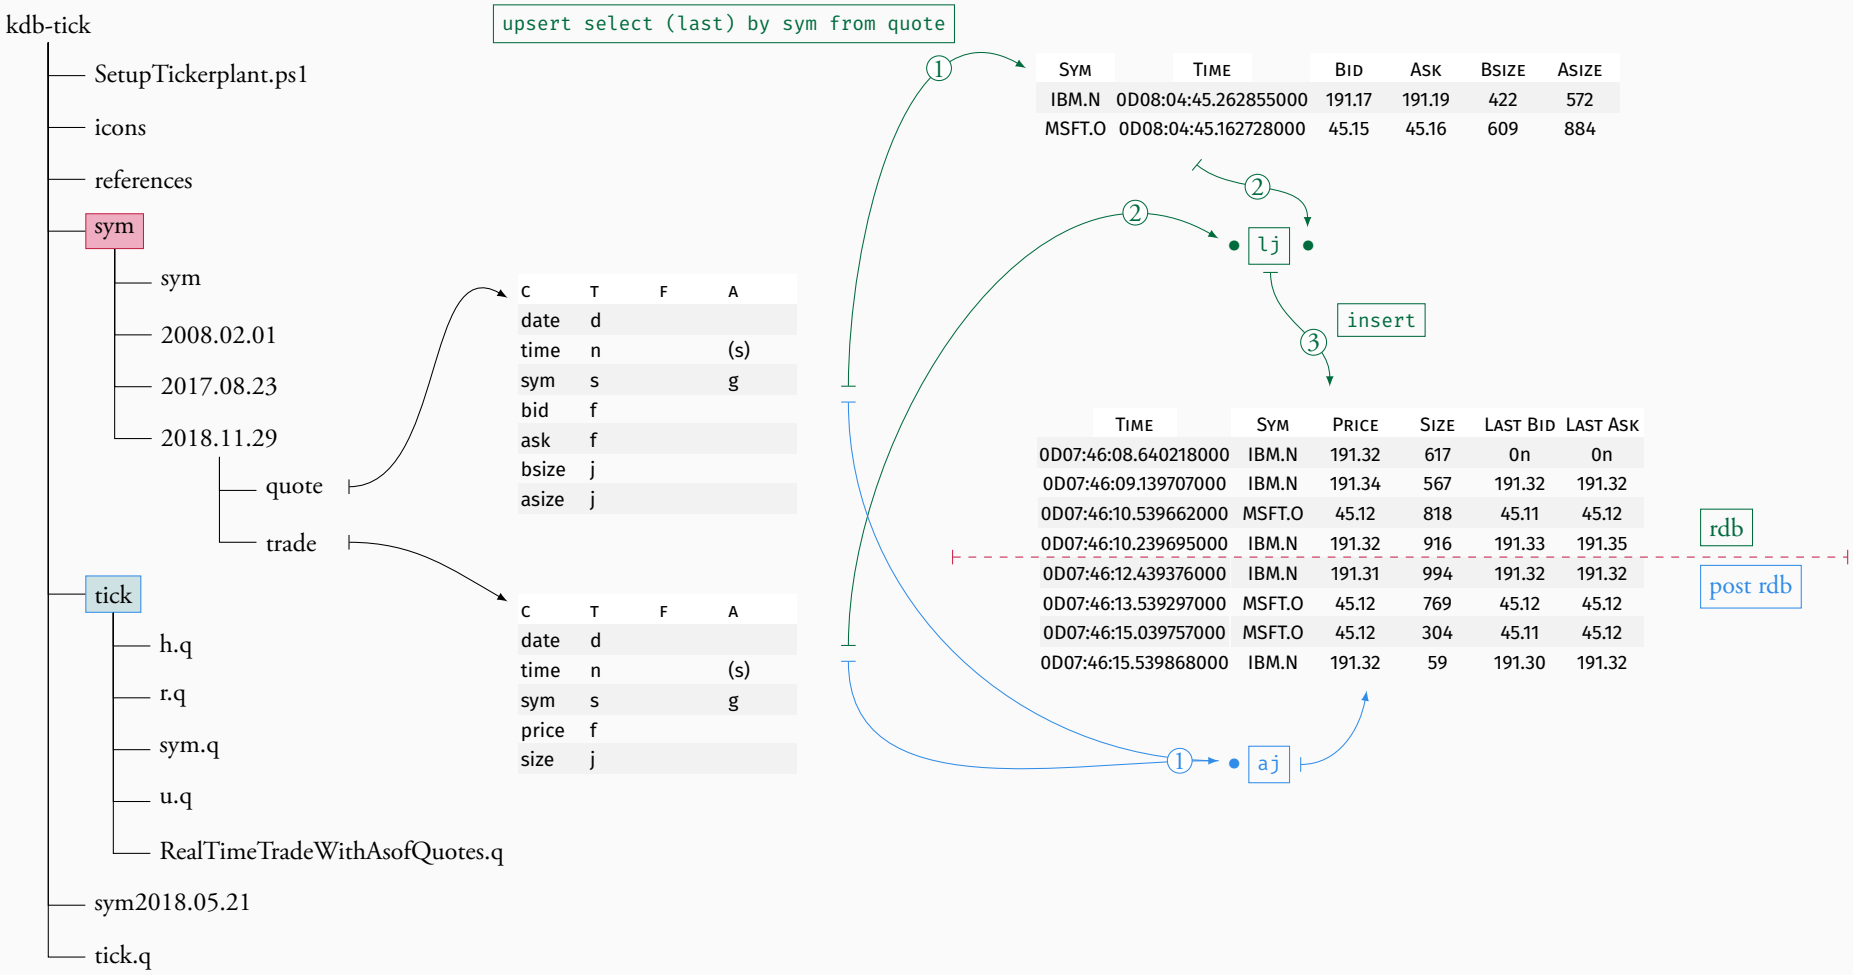

Temporal Joins

function	description
aj/aj0	get the value of a field in one table asof the value in another
asof	like aj except can be passed a dictionary as left argument of the sym/times to return data asof
wj	generalization of aj with an additional argument list of pairs of start/end times that specify windows
raj	reverse aj for most recent future looking data points; simplest implementation negate the date

TEMPORAL ARITHMETIC

function	description
date ± integers	integers representing day counts can be added/subtracted from dates and as a result dates themselves can be added to/subtracted from to give day counts
time ± integers	integers representing milliseconds can be added/subtracted from times
datetime ± floats	the integer part represents a day count and the fractional part represents a fraction of a day `2003.03.23T08:31:53.000 + 2.5%24`
timestamp ± floats	does not work
timestamp ± integers	integers representing milliseconds can be added/subtracted
timestamp ± 1D	integers with trailing ”D” representing days can be added/subtracted

timestamp × integers	multiply minutes/hours/days etc by integer
	`0D00:12*3`
	`0D01:00*2`
	`2012.02.28D00:18:00+(0D00:10+til 115)`

.Q.addmonths	`.Q.addmonths[.z.P;12]`
addyear	`addyear[.z.P;1]`

EXAMPLE: AJ VERSUS WJ

Listing 1: Q/KDB+ code for aj vs. wj

```
t:([]sym:3#`ibm;time:10:01:01 10:01:04 10:01:08;price:100 101 105)
q:([]sym:9#`ibm;time:10:01:01+til 9;ask:101 103 103 104 104 107 108 107 108;
     bid:98 99 102 103 103 104 106 106 107)
w:-2 1+\:t.time
w:0 10+\:t.time
wj[w;`sym`time;t;(q;(max;`ask);(min;`bid))]
aj[`sym`time;t;q]
```

	asks and bids			trades		start/end times of bins		wj		aj	
sym	time	ask	bid	time	price	w1	w2	ask	bid	ask	bid
‘ibm	10:01:01	101	98	10:01:01	100	10:01:01	10:01:11	108	98	101	98
‘ibm	10:01:02	103	99								
‘ibm	10:01:03	103	102								
‘ibm	10:01:04	104	103	10:01:04	101	10:01:04	10:01:14	108	103	104	103
‘ibm	10:01:05	104	103								
‘ibm	10:01:06	107	104								
‘ibm	10:01:07	108	106								
‘ibm	10:01:08	107	106	10:01:08	105	10:01:08	10:01:18	108	106	107	106
‘ibm	10:01:09	108	107								
‘ibm	10:01:10	108	107								

Architecture

Single Console Tickerplant

	process	description
	Tickerplant	writes to logfile in addition to publishing data to RDB and other subscribers
	Feedhandler	connects, recieves & normalizes exchange/market data before feeding it to the tickerplant
	RDB	real-time in-memory store for current day's data
	HDB	partitioned on-disk stores for historical data
	Logfile	on-disk file that stores alll intra-day updates for potential re-play during recovery
	Chained TP	subscribes to main (zero-latency) tickerplant and recieves de-layed updates and keeps no log file
	Chained RDB	subscribes to either the main or chained tickerplant

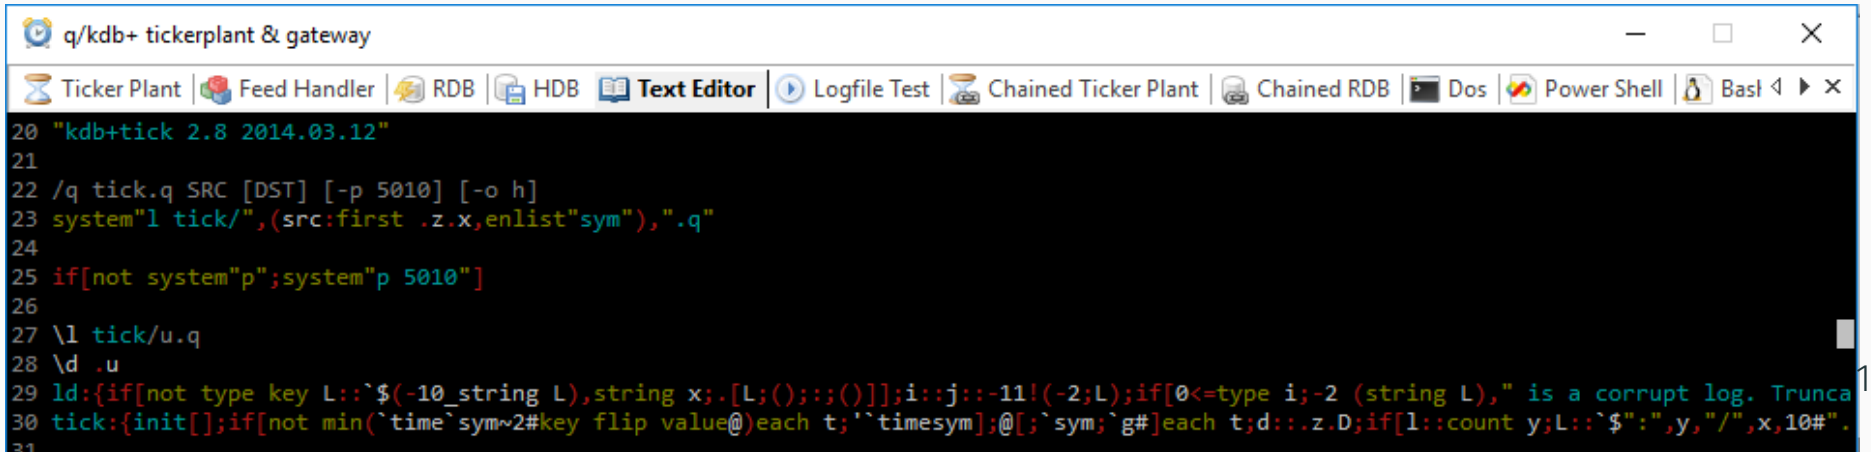

Data Layer

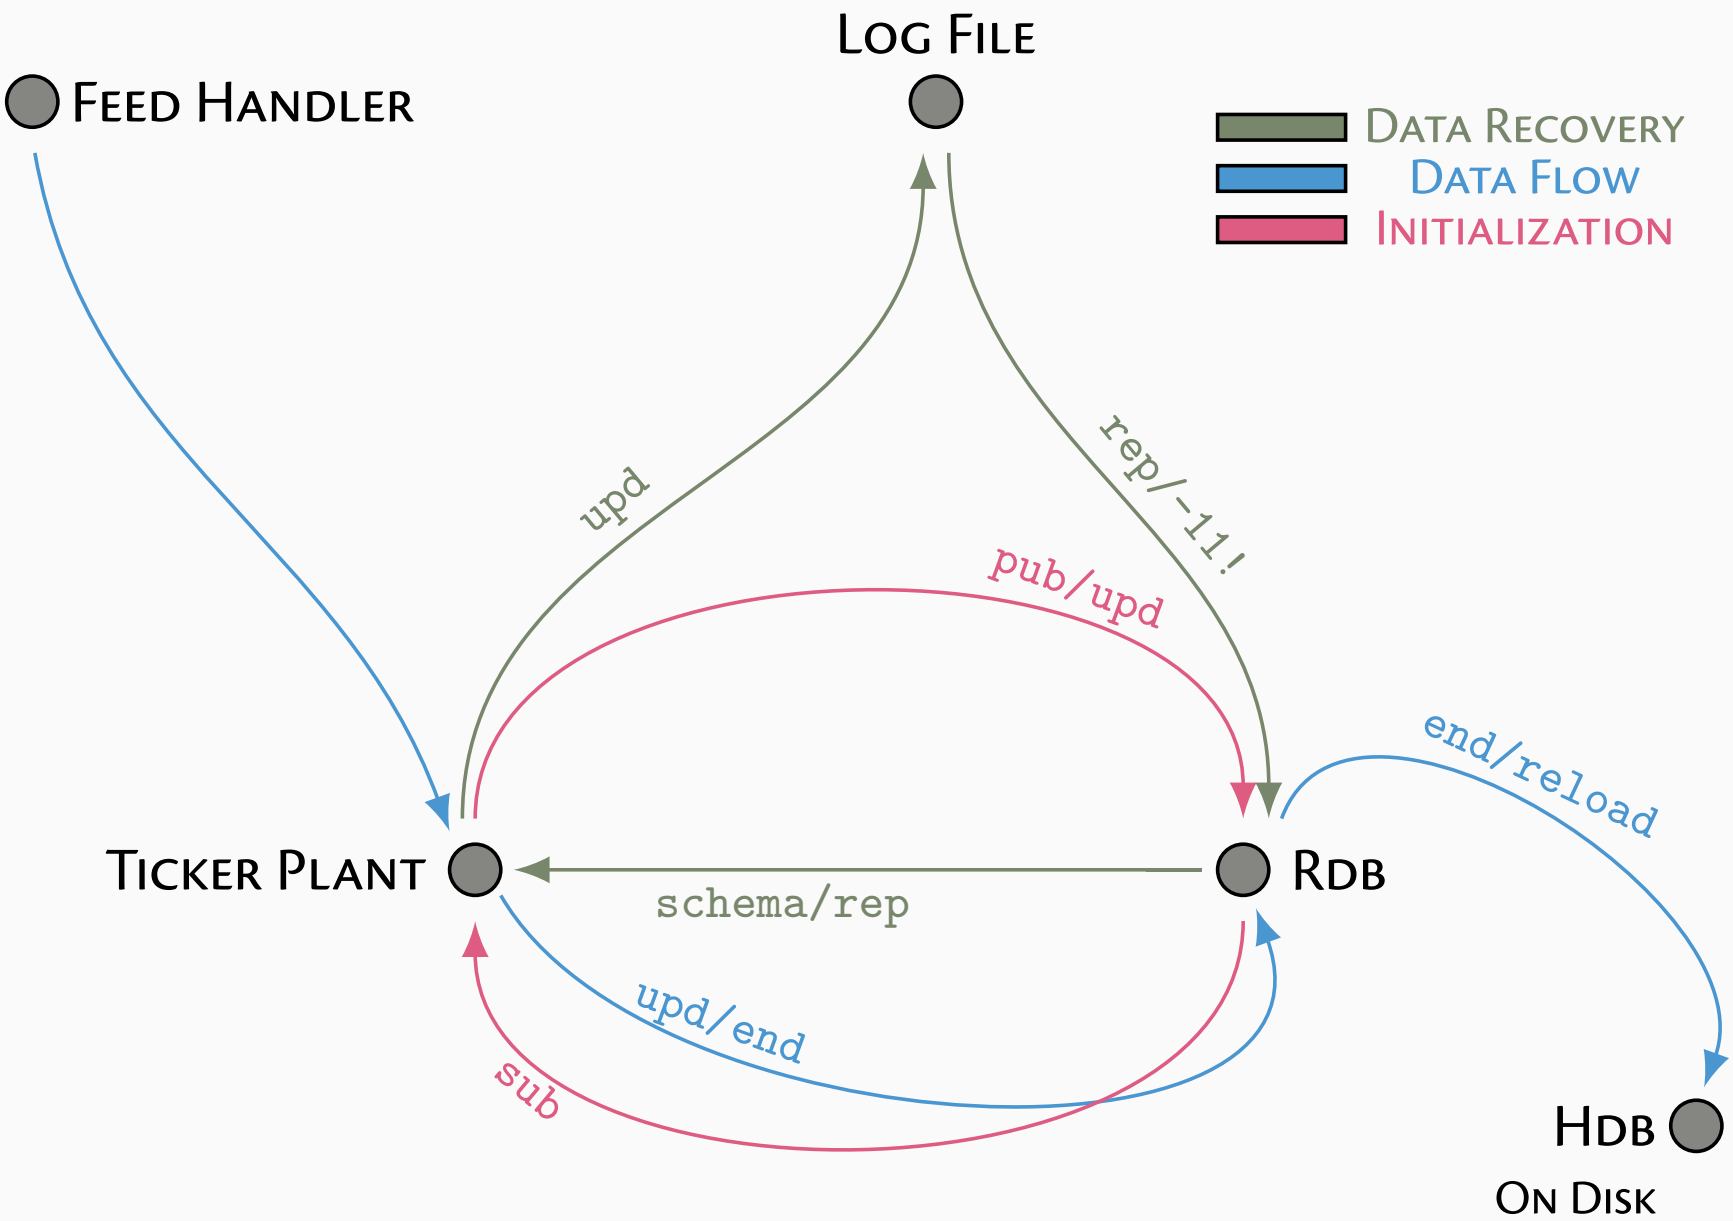
Log File
Feed Handler
Data Recovery
Data Flow
Initialization
upd
rep/-11!
pub/upd
end/reload
Ticker Plant
schema/rep
Rdb
upd/end
sub
Hdb
On Disk

Data Access Layer

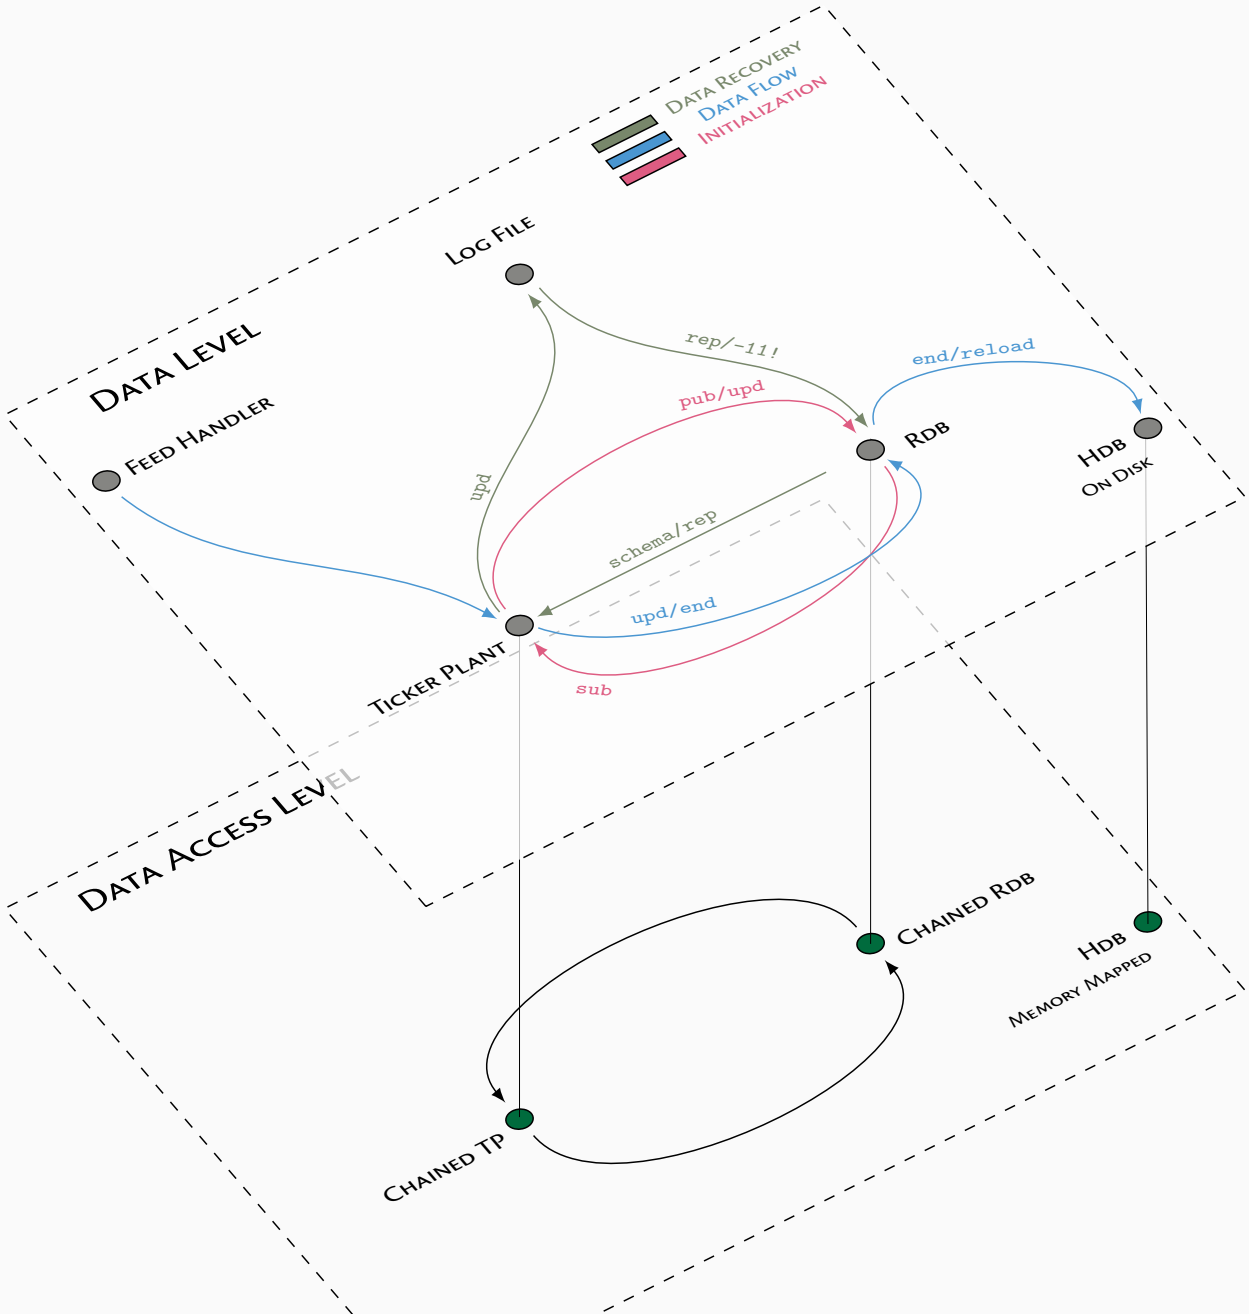

Application Layer

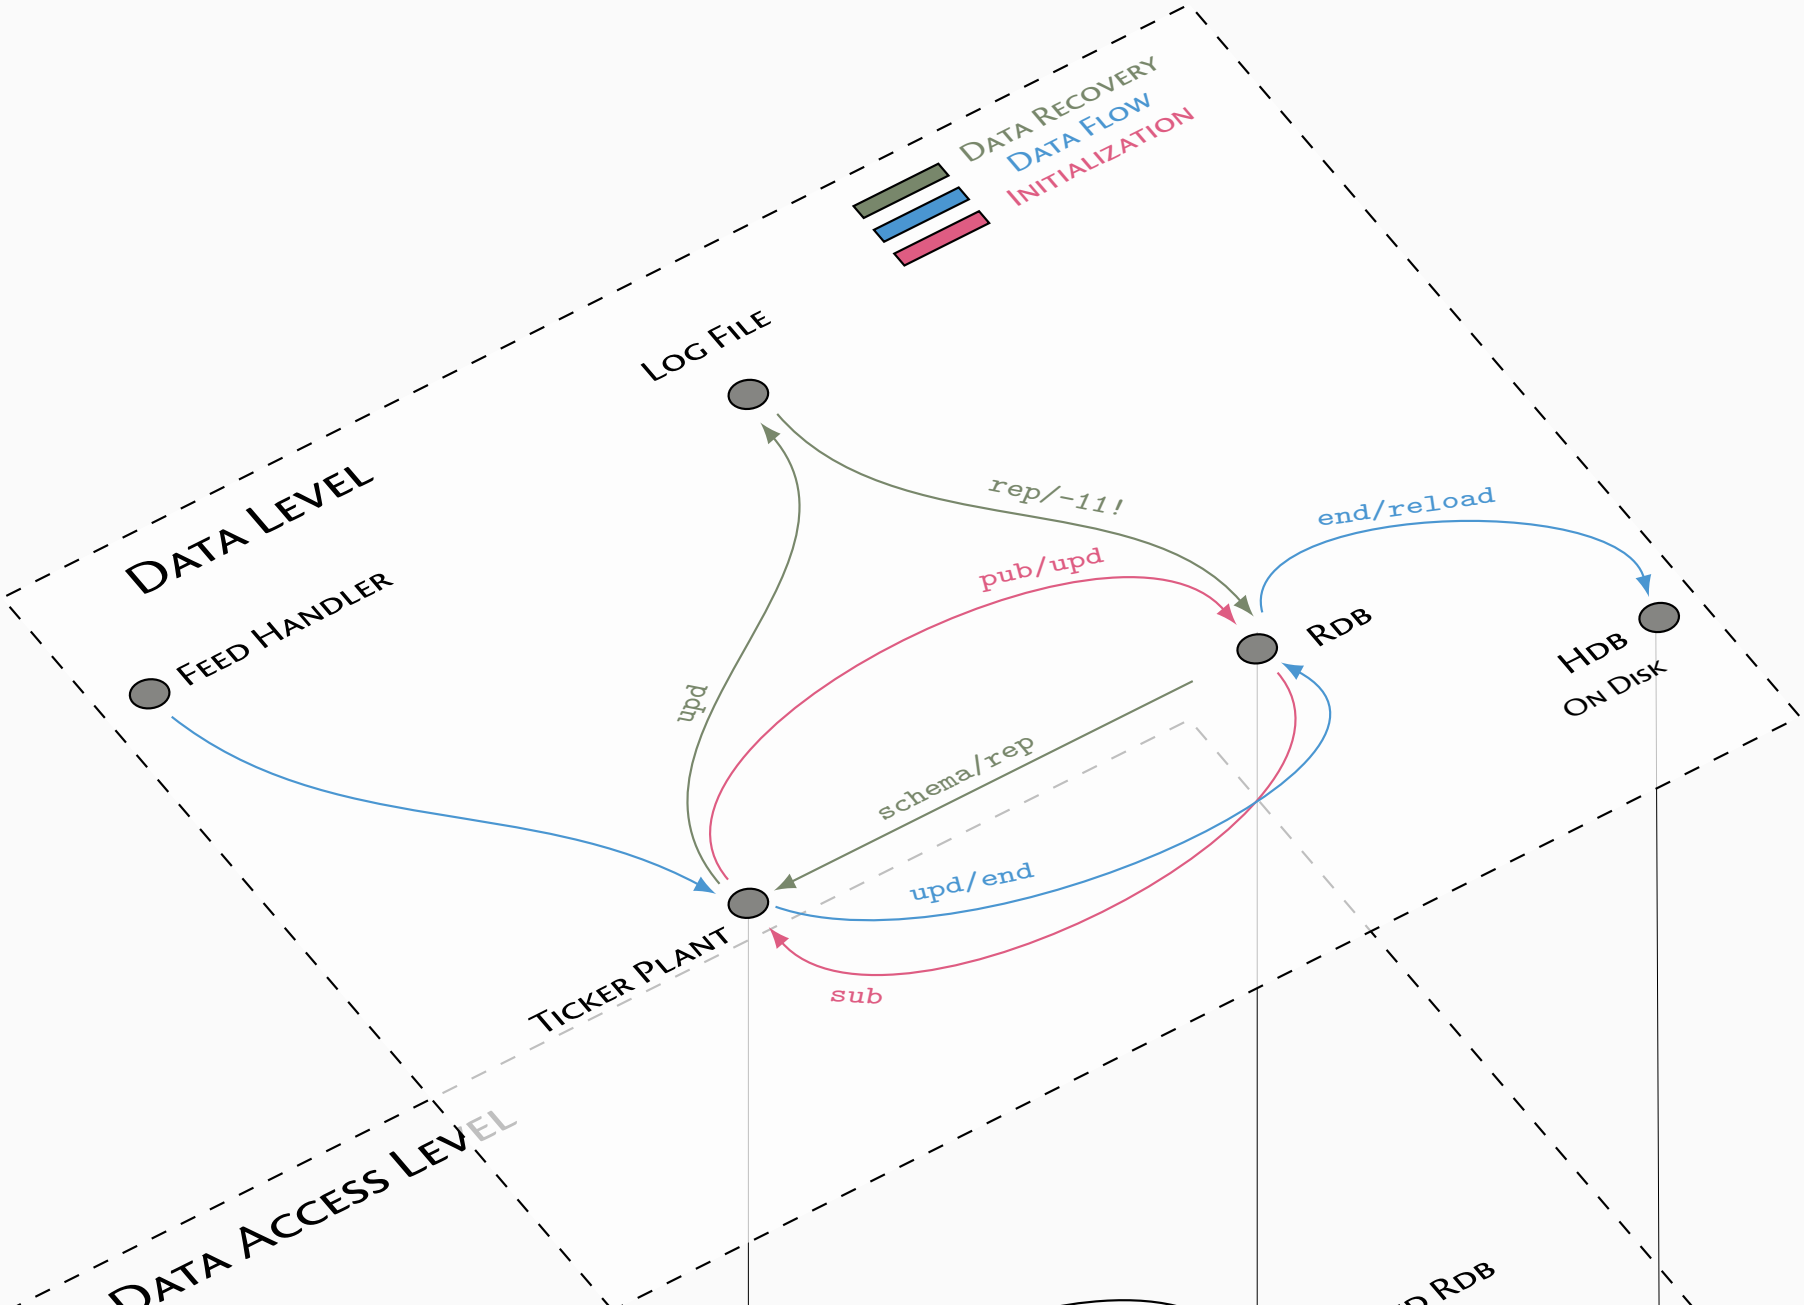

Single Console Tickerplant Screenshot

Installation

Install Using Powershell

Download powershell script from :

https://github.com/rohanshiloh/tickerplant.ai

ps1 script

- Downloads 64-bt console2 and extracts it
- Downloads all required q files from code.kx.com
- Downloads all reference white-papers from code.kx.com
- Copies console2 config file into appropriate folder
- Checks which verson of python and powershell are installed
- Starts single console tickerplant

Setup Python, Bash & VC Compiler

PYTHON SETUP

- Install Anaconda and iPython
- Identify your python path and add it to the PATH environment
- follow instructions for experimental setup of pyQ here `https://pyq.enlnt.com/install/install.html#experimental-support-for-windows`
- Set console2's shell value :

```
cmd.exe /k C:\Anaconda\IPython.exe
```

BASH SETUP

- Install the Linux Subsystem on Windows
- Set console2's shell value :

```
%SystemRoot%\system32\bash.exe
```

Conclusion

Summary

The application of REAL-TIME ITERATIVE machine learning algorithms to stock market order book data can most *efficiently* be done using the Q/KDB+/PYQ/PYTHON stack.

`github.com/rohanshiloh/tickerplant.ai`

Next Steps

- Tick PCA & Data Visualization of the Order Book at Depth
- Frequentist, Bayesian, Generalized & Regularized (Ridge L_2 & Lasso L_1) Linear Regression Models
- Bayesian (Multi-Strategy & Multi-Period) Portfolio Optimization
- Structural (Regularized) Covariance

References i

N. Coulter, *Order book: a kdb+ intraday storage and access methodology*, 2014.

Coulter covers the fine balance between processing/storage of an order book and it's subsequent retrieval and analysis. Different trade table schemas and accessor functions are presented and dicussed.

D. Easley and M. O'Hara, *Price, trade, size and information in securities markets*, 1987.

References ii

C. M. C. Lee and M. J. Ready, *Inferring trade direction from intraday data*, 1991.

Lee and Ready's seminal paper on tick data anaysis came out in 1991. They did their research on a supercomputer at Cornell; likely not running q/kdb. Oddly enough Canada supported their research. The paper outlines various 'tick tests' for an uptick, downtick, a zero-uptick and a zero-downtick. The additional function tradeclass in the file tickaddendum.q provides defintions for each of these classes and appends them as a column to the trade table.

C. McCarthy, *Intraday writedown solutions*, 2014.

This paper will be incredibly useful to independent researchers running a tickerplant at home on 4core/4gb.

References iii

N. Perrem, *Building real-time tick subscribers*, 2014.
Perrem's excellent paper is a good introduction to standard ticker plant architecture. Chapter 4 covers how to build a custom temporally joined trade aj quote architecture with it's own binary logfile.